


Branch: master

econ128 / Labs / Lab1 / Lab1_PartC_Exploratory_Data_Analysis.ipynb

Find file

Copy path

 msalloum adding labs 1, 4, 5, 6 d160664 on Aug 8

1 contributor

1159 lines (1158 sloc) 76.4 KB

Exploratory Data Analysis

The HOUSES dataset contains a collection of recent real estate listings in San Luis Obispo county and around it. The dataset is as a CSV file. The dataset contains the following fields:

1. MLS: Multiple listing service number for the house (unique ID).
2. Location: city/town where the house is located. Most locations are in San Luis Obispo county and northern Santa Barbara county (Santa Maria-Orcutt, Lompoc, Guadalupe, Los Alamos), but there some out of area locations as well.
3. Price: the most recent listing price of the house (in dollars).
4. Bedrooms: number of bedrooms.
5. Bathrooms: number of bathrooms.
6. Size: size of the house in square feet.
7. Price/SQ.ft: price of the house per square foot.
8. Status: type of sale. Thee types are represented in the dataset: Short Sale, Foreclosure and Regular.

Lets import the required libraries that we will be using later.

```
In [2]: from numpy import *
import pandas as pd
```

Let's load the dataset into a pandas dataframe and have a look at the headers.

```
In [3]: df = pd.read_csv('data.csv', sep=',', error_bad_lines=False) # read fie as a dataframe

print len(df)      # print number of rows
print list(df)      # print 1st ten headers

781
['MLS', 'Location', 'Price', 'Bedrooms', 'Bathrooms', 'Size', 'Price/SQ.Ft', 'Status']
```

Lets take a look at the first 2 rows of the dataframe.

```
In [17]: df.head(2)
```

```
Out[17]:
```

	MLS	Location	Price	Bedrooms	Bathrooms	Size	Price/SQ.Ft	Status
0	132842	Arroyo Grande	795000	3	3	2371	335.30	Short Sale
1	134364	Paso Robles	399000	4	3	2818	141.59	Short Sale

Next, lets look at a specific column or feature in the dataframe.

```
In [4]: print "Bedrooms:", df.Bedrooms.unique()
print "Bathrooms:", df.Bathrooms.unique()

Bedrooms: [ 3  4  2  7  1  5  0  6 10]
Bathrooms: [ 3  4  1  2  5  7  6 11]
```

What if we want to drop a column from the dataframe?

```
In [119]: df = df.drop('Location', 1)
```

Let's rename the first column.

Hint: A Google search for 'python pandas dataframe rename' points you at this documentation (<http://pandas-docs.github.io/pandas-docs-travis/>).

```
In [18]: print df.columns
df.rename(columns={df.columns[0]:'listing id'}, inplace=True)
df.columns

Index([u'MLS', u'Location', u'Price', u'Bedrooms', u'Bathrooms', u'Size',
       u'Price/SQ.Ft', u'Status'],
      dtype='object')

Out[18]: Index([u'listing id', u'Location', u'Price', u'Bedrooms', u'Bathrooms',
               u'Size', u'Price/SQ.Ft', u'Status'],
              dtype='object')
```

What is the max, min, mean/avg, and standard deviation of the column 'bedroomcnt'?

```
In [5]: print "Min: " , min(df['Price/SQ.Ft'])
print "Max: " , max(df['Price/SQ.Ft'])
print "Std: " , df['Price/SQ.Ft'].std()
print "Avg: " , df['Price/SQ.Ft'].mean()
```


```
Min: 19.33
Max: 1144.64
Std: 115.082145976
Avg: 213.131293214
```

Plot the distribution of 'yearbuild' using matplotlib

```
In [13]: import matplotlib.mlab as mlab
import matplotlib.pyplot as plt
import numpy as np

# remove NAN from array
x = df['Price/SQ.Ft'][~np.isnan(df['Price/SQ.Ft'])]

# plot histogram
n, bins, patches = plt.hist(x, 30, normed=1, facecolor='blue')
plt.show()
```




One of the best ways to inspect data is visualize it. One way to do this is by using a scatter plot. A scatter plot of the data puts one feature along the x-axis and another along the y-axis, and draws a dot for each data point.

Since its difficult to visualize more than 2 or 3 features, one possibility is to use a pair plot that looks at all possible pairs of features. The pair plot shows the interaction of each pair of features inorder to visualize any correlation between features.

```
In [15]: # import the scatter_matrix functionality
import random as rand
from pandas.tools.plotting import scatter_matrix
import matplotlib.pyplot as plt

print df.shape
x = df.iloc[:,[1,2,3,4,5]] # extract only 5 columns from dataframe (using index)
y = x.dropna(thresh=2) # drop any rows that have 2 or more fields as NAN
a = pd.scatter_matrix(x, alpha=0.05, figsize=(5,5), diagonal='hist')
plt.show()
```

(781, 8)



Price 1e7 Bedrooms Bathrooms Size

Let's create a crosstabulation or contingency table of the factors.

Hint: A Google search for 'python pandas cross tabulation' points you at this documentation (<http://pandas.pydata.org/pandas-docs/stable/reshaping.html#cross-tabulations>).

```
In [16]: pd.crosstab(df.Bedrooms, df.Bathrooms)
```

Out[16]:

Bathrooms	1	2	3	4	5	6	7	11
Bedrooms								
0	2	0	0	0	0	0	0	0
1	10	1	0	0	0	0	0	0
2	33	78	12	0	0	0	0	0
3	20	286	119	5	1	0	0	0
4	0	76	76	19	6	0	0	0
5	0	3	16	9	0	2	1	0
6	0	0	3	0	0	0	1	0
7	0	0	0	0	1	0	0	0
10	0	0	0	0	0	0	0	1

```
In [37]: fig=plt.figure()
plt.scatter(df.Price, df.Size)
axis = fig.gca() #get current axis
axis.set_title('Price vs Size')
axis.set_xlabel('Price')
axis.set_ylabel('Size')
fig.canvas.draw()
```

Pandas Data Munging

The first concept we deal with here is pandas groupby. The idea is to group a dataframe by the values of a particular factor variable. The documentation can be found here (<http://pandas-docs.github.io/pandas-docs-travis/>).

```
In [43]: status_groupby = df.groupby('Status')
status_groupby.head()
```

Out[43]:

	listing id	Location	Price	Bedrooms	Bathrooms	Size	Price/SQ.Ft	Status
0	132842	Arroyo Grande	795000	3	3	2371	335.30	Short Sale
1	134364	Paso Robles	399000	4	3	2818	141.59	Short Sale
2	135141	Paso Robles	545000	4	3	3032	179.75	Short Sale
3	135712	Morro Bay	909000	4	4	3540	256.78	Short Sale
4	136282	Santa Maria-Orcutt	109900	3	1	1249	87.99	Short Sale
53	143436	Templeton	1399000	4	3	6500	215.23	Foreclosure
56	143534	Morro Bay	789000	3	3	2100	375.71	Foreclosure
63	144314	Morro Bay	899000	3	3	2430	369.96	Foreclosure
64	144316	Morro Bay	1045000	3	3	2100	497.62	Foreclosure
65	144318	Morro Bay	774000	2	2	1550	499.35	Foreclosure
617	154325	Lompoc	149900	3	1	1000	149.90	Regular
621	154329	Santa Maria-Orcutt	177000	3	2	1500	118.00	Regular
622	154330	Nipomo	122500	4	2	1248	98.16	Regular
626	154343	Arroyo Grande	425000	3	2	1437	295.76	Regular
628	154345	Nipomo	175000	2	2	1152	151.91	Regular

The function `groupby` gives you a dictionary-like object, with the keys being the values of the factor, and the values being the corresponding subsets of the dataframe.

```
In [47]: # lets print the number of rows per each type of status
for key, value in status_groupby:
    print "( key, len(value) ) = ( ", key, ", ", len(value), " )"
    v=value

( key, len(value) ) = ( Foreclosure , 162 )
( key, len(value) ) = ( Regular , 103 )
( key, len(value) ) = ( Short Sale , 516 )
```

The `groupby` function also acts like an object that can be mapped. After the mapping is complete, the rows are put together (reduced) into a larger dataframe. For example, using the `describe` function. The documentation of the `describe` function can be found here (<http://pandas-docs.github.io/pandas-docs-travis/>).

```
In [52]: status_df =status_groupby['Bathrooms','Bedrooms','Price','Size','Price/SQ.Ft'].describe()
print type(status_df)
status_df.head(20)
```

```
<class 'pandas.core.frame.DataFrame'>
```

Out[52]:

		Bathrooms	Bedrooms	Price	Size	Price/SQ.Ft
Status						
Foreclosure	count	162.000000	162.000000	1.620000e+02	162.000000	162.000000
	mean	2.314815	3.216049	3.632146e+05	1762.493827	200.221605
	std	0.807359	0.956956	2.701240e+05	880.295908	97.433947
	min	1.000000	0.000000	2.900000e+04	628.000000	19.330000
	25%	2.000000	3.000000	1.811250e+05	1200.000000	140.050000
	50%	2.000000	3.000000	2.789500e+05	1500.000000	186.780000
	75%	3.000000	4.000000	4.437250e+05	2064.000000	225.525000
	max	7.000000	6.000000	1.799000e+06	6500.000000	686.020000
Regular	count	103.000000	103.000000	1.030000e+02	103.000000	103.000000
	mean	2.446602	3.048544	6.411621e+05	1926.699029	311.406796
	std	0.801080	0.832886	6.752140e+05	826.470618	189.019573
	min	1.000000	1.000000	2.650000e+04	120.000000	19.720000
	25%	2.000000	3.000000	3.154500e+05	1431.000000	200.100000
	50%	2.000000	3.000000	5.090000e+05	1832.000000	282.690000
	75%	3.000000	3.000000	7.620000e+05	2237.500000	369.740000
	max	6.000000	6.000000	5.499000e+06	5060.000000	1144.640000
Short Sale	count	516.000000	516.000000	5.160000e+02	516.000000	516.000000
	mean	2.350775	3.137597	3.381775e+05	1718.463178	197.567326
	std	0.867436	0.825834	2.365197e+05	795.330609	88.248737
	min	1.000000	0.000000	4.090000e+04	398.000000	34.050000

Another way of printing statistics using `group by` is shown below, without having to iterate through each item in the dataframe. You can find documentation about `std` function here (<http://pandas-docs.github.io/pandas-docs-travis/>).

```
In [57]: status_df =df.groupby('Status').std()
print status_df
```

```

           listing id      Price  Bedrooms  Bathrooms      Size \
Status
Foreclosure  2691.998699  270124.016966  0.956956    0.807359  880.295908
```